

Shock Agent Suite — ADK/MCP Repo Skeleton (Google AI Agent Intensive)

This repository blueprint implements four advanced agents with on-demand tool synthesis, causal/chronosynaptic research, machine-checkable reasoning, and world-simulation planning. It is aligned with ADK function tools and MCP exposure. Copy this structure into a new repo and fill in TODOs.

0) Repository Layout

```
shock-agent-suite/
├── README.md
├── .env.example
├── .gitignore
├── docker/
│   └── Dockerfile
├── scripts/
│   ├── dev_up.sh
│   ├── run_demo.sh
│   └── seed_fixtures.sh
├── adk_app/
│   ├── app.py
│   └── tools/
│       ├── web_fetch.py
│       ├── html_extract.py
│       ├── pdf_tables.py
│       ├── notes_graph_add.py
│       ├── tool_specify.py
│       ├── tool_codegen.py
│       ├── tool_sandbox_run.py
│       ├── tool_register.py
│       ├── causal_tools.py
│       ├── tda_tools.py
│       ├── smt_tools.py
│       ├── fuzz_tools.py
│       ├── world_model_tools.py
│       └── guard_tools.py
└── policies/
    ├── tool_allowlist.yaml
    └── safety_rules.md
└── mcp_server/
    └── server.py
```

```
|   |   └─ schemas/
|   |       └─ web_fetch.json
|   |       └─ html_extract.json
|   |       └─ pdf_tables.json
|   |       └─ notes_graph_add.json
|   |       └─ tool_specify.json
|   |       └─ tool_codegen.json
|   |       └─ tool_sandbox_run.json
|   |       └─ tool_register.json
|   |       └─ causal_tools.json
|   |       └─ smt_tools.json
|   |       └─ world_model_tools.json
|   └─ registry/
|       └─ README.md
└─ sandbox/
    └─ worker.py
    └─ policy.py
    └─ tests/
        └─ sample_oracles.json
        └─ tool_examples/
            └─ pdf_ocr_tables/
                └─ fixture.pdf
                └─ expected.json
└─ engines/
    └─ chronosynaptic/
        └─ graph_store.py
        └─ episodes.py
        └─ planner.py
    └─ program_of_thought/
        └─ ir.py
        └─ compiler.py
        └─ smt_adapter.py
        └─ viz.py
    └─ world_sim/
        └─ model.py
        └─ rollout.py
        └─ evc.py
    └─ tests/
        └─ test_tool_foundry.py
        └─ test_causal_planner.py
        └─ test_pot_auditor.py
        └─ test_world_sim.py
```

1) Quickstart

```
# 1) Clone and setup  
cp .env.example .env  
python -m venv .venv && source .venv/bin/activate  
pip install -r requirements.txt  
  
# 2) Start ADK app + MCP server (dev)  
../scripts/dev_up.sh  
  
# 3) Seed fixtures for demos  
../scripts/seed_fixtures.sh  
  
# 4) Run the 90s Tool-Foundry demo  
../scripts/run_demo.sh tool-foundry
```

requirements.txt (excerpt):

```
fastapi  
uvicorn  
httpx  
lxml  
beautifulsoup4  
pydantic  
networkx  
scikit-learn  
numpy  
scipy  
pandas  
pillow  
pytesseract  
pdfminer.six  
camelot-py[cv]  
python-multipart  
z3-solver  
hypothesis  
matplotlib  
orjson  
pyyaml
```

2) ADK Agent Host (adk_app/app.py)

```
from google.adk.agents import LlmAgent
from google.adk.tools import Tool
from google.genai.types import Tool as GeminiTool, UrlContext

from tools.web_fetch import web_fetch
from tools.html_extract import html_extract
from tools.pdf_tables import pdf_tables
from tools.notes_graph_add import notes_graph_add
from tools.tool_specify import tool_specify
from tools.tool_codegen import tool_codegen
from tools.tool_sandbox_run import tool_sandbox_run
from tools.tool_register import tool_register
from tools.causal_tools import corpus_ingest, causal_propose, causal_intervene,
evidence_verify
from tools.tda_tools import tda_summarize
from tools.smt_tools import pot_compile, smt_check, proof_export
from tools.fuzz_tools import fuzz_casegen
from tools.world_model_tools import sim_train, sim_rollout, act_execute
from tools.guard_tools import guard_estimate_risk, rollback_recover

agent = LlmAgent(
    model="gemini-2.5-flash",
    name="shock_agent_suite",
    description="Self-extending tool foundry + causal planner + PoT auditor + world-sim planner.",
    instruction=(
        "Prefer url_context for retrieval. For new capabilities, synthesize a tool via"
        " tool_specify->tool_codegen->tool_sandbox_run->tool_register before use."
        " For research, maintain claims in notes_graph_add."
        " For reasoning, compile to PoT IR and check with smt_check; export proof."
        " For projects, simulate with sim_rollout under risk/cost guardrails."
    ),
    tools=[
        Tool(fn=web_fetch), Tool(fn=html_extract), Tool(fn=pdf_tables),
        Tool(fn=notes_graph_add),
        Tool(fn=tool_specify), Tool(fn=tool_codegen), Tool(fn=tool_sandbox_run),
        Tool(fn=tool_register),
        Tool(fn=corpus_ingest), Tool(fn=tda_summarize), Tool(fn=causal_propose),
        Tool(fn=causal_intervene), Tool(fn=evidence_verify),
        Tool(fn=pot_compile), Tool(fn=smt_check), Tool(fn=proof_export),
        Tool(fn=fuzz_casegen),
        Tool(fn=sim_train), Tool(fn=sim_rollout), Tool(fn=act_execute),
    ]
)
```

```

    Tool(fn=guard_estimate_risk), Tool(fn=rollback_recover),
        GeminiTool(url_context=UrlContext),
    ],
)

```

3) Core Drop-in Tools (summaries + TODOs)

3.1 web_fetch.py

```

import httpx, time
from bs4 import BeautifulSoup

def web_fetch(url: str, timeout_sec: int = 20) -> dict:
    try:
        with httpx.Client(follow_redirects=True, timeout=timeout_sec) as s:
            r = s.get(url)
            ct = (r.headers.get("content-type") or "").lower()
            text = r.text
            title, links, md = None, [], text
            if "html" in ct or text.strip().startswith("<"):
                soup = BeautifulSoup(text, "lxml")
                title = soup.title.string if soup.title else None
                for t in soup(["script", "style", "noscript"]): t.decompose()
                links = [a.get("href") for a in soup.find_all("a") if a.get("href")]
                md = "\n".join(h.get_text(" ", strip=True) for h in
soup.find_all(["h1", "h2", "h3", "p", "li"]))
                return {"title": title, "markdown": md, "links": links[:200],
"content_type": ct}
            except Exception as e:
                return {"isError": True, "message": str(e)}

```

3.2 html_extract.py

```

from bs4 import BeautifulSoup
import re

def html_extract(markup: str, selector: str, attrs: list[str] | None = None) ->
dict:
    soup = BeautifulSoup(markup, "lxml")
    if selector.startswith("table"):
        m = re.match(r"table(?:\[([^\d+)\])?$", selector)
        idx = int(m.group(1)) if m and m.group(1) else 0
        tables = soup.find_all("table")

```

```

    if idx >= len(tables):
        return {"rows": []}
    rows = []
    for tr in tables[idx].find_all("tr"):
        cells = [c.get_text(" ", strip=True) for c in
tr.find_all(["th","td"])]
        if cells: rows.append(cells)
    return {"rows": rows}
nodes = soup.select(selector)
out = []
for n in nodes:
    item = {"text": n.get_text(" ", strip=True)}
    if attrs: item["attrs"] = {a: n.get(a) for a in attrs}
    out.append(item)
return {"matches": out}

```

3.3 pdf_tables.py (skeleton)

```

import subprocess, tempfile, json, os
# TODO: implement with camelot/tabula fallback + OCR (pytesseract) for scanned
PDFs

def pdf_tables(url: str, pages: str | None = None) -> dict:
    # Download to temp; try camelot; fallback to OCR grid extraction
    return {"tables": [], "note": "TODO: implement PDF table extraction
pipeline"}

```

3.4 notes_graph_add.py

```

import time, uuid, networkx as nx
G = nx.MultiDiGraph()

def notes_graph_add(nodes: list[dict], edges: list[dict] | None = None) -> dict:
    ids = []
    for n in nodes:
        nid = n.get("id") or str(uuid.uuid4())
        G.add_node(nid, **{k:v for k,v in n.items() if k != 'id'})
        ids.append(nid)
    for e in edges or []:
        G.add_edge(e["src"], e["dst"], **{k:v for k,v in e.items() if k not in
('src','dst')})
    return {"node_ids": ids, "node_count": G.number_of_nodes(), "edge_count": G.number_of_edges()}

```

4) Tool-Foundry (Create/Test/Register Tools On-the-Fly)

4.1 tool_specify.py

```
from typing import Dict

def tool_specify(task_text: str) -> Dict:
    # Heuristic sketch: derive name/description/input_schema and 2-3 oracle
    tests
    # TODO: replace with LLM-backed spec generator constrained by JSON schema
    return {
        "name": "pdf_ocr_tables",
        "description": "Extract tables from scanned PDFs using OCR then
structure rows.",
        "input_schema": {
            "type": "object",
            "properties": {"url": {"type": "string"}, "pages": {"type": [
                "string", "null"]}},
            "required": ["url"]
        },
        "oracle_tests": [
            {"input": {"url": "fixtures/fixture.pdf"}, "expect": {"tables_min": 1}}
        ]
    }
```

4.2 tool_codegen.py

```
from typing import Dict

def tool_codegen(spec: Dict) -> Dict:
    # Emit a Python tool file from the spec (simple template)
    name = spec["name"]
    code = """
import json
from typing import Dict

def {name}(url: str, pages: str | None = None) -> Dict:
    # TODO: real OCR + table extractor; this is a stub
    return {"tables": [], "note": "stub"}
"""
    return {"files": [{"path": f"generated/{name}.py", "content": code}],
    "entrypoint": f"generated/{name}.py", "build_cmd": ""}
```

4.3 tool_sandbox_run.py

```
import time, resource, subprocess, tempfile, os, json

def tool_sandbox_run(files: list[dict], cmd: str = "python -c 'print(0)'",
tests: list[dict] | None = None, limits: dict | None = None) -> dict:
    limits = limits or {"cpu_sec": 5, "mem_mb": 512, "net": False}
    with tempfile.TemporaryDirectory() as td:
        for f in files:
            p = os.path.join(td, f["path"])
            os.makedirs(os.path.dirname(p), exist_ok=True)
            with open(p, "w", encoding="utf-8") as h: h.write(f["content"])
        start = time.time()
        try:
            proc = subprocess.run(cmd, cwd=td, shell=True, capture_output=True,
text=True, timeout=limits["cpu_sec"])
            wall_ms = int((time.time()-start)*1000)
        except subprocess.TimeoutExpired:
            return {"pass": False, "failures": [{"message": "timeout"}],
"metrics": {"wall_ms": limits["cpu_sec"]*1000}}
        failures = []
        # TODO: import generated tool and run tests as callables
        return {"pass": len(failures)==0, "failures": failures, "metrics":
{"wall_ms": wall_ms}}
```

4.4 tool_register.py

```
import json, os
REGISTRY = "mcp_server/registry"

def tool_register(schema: dict) -> dict:
    os.makedirs(REGISTRY, exist_ok=True)
    name = schema.get("name")
    path = os.path.join(REGISTRY, f"{name}.json")
    with open(path, "w", encoding="utf-8") as h:
        json.dump(schema, h, indent=2)
    return {"tool_id": name, "version": schema.get("version", "0.1.0"), "path":
path}
```

5) Causal / Chronosynaptic Research Tools

`engines/chronosynaptic/graph_store.py`

```
import networkx as nx
class GraphStore:
    def __init__(self):
        self.G = nx.MultiDiGraph()
    def add_claim(self, text, source_id, confidence: float):
        nid = f"claim:{hash(text)}"
        self.G.add_node(nid, type="claim", text=text, confidence=confidence,
source_id=source_id)
        return nid
    def link(self, a, b, etype, **props):
        self.G.add_edge(a, b, type=etype, **props)
```

`adk_app/tools/causal_tools.py (stubs)`

```
def corpus_ingest(uri: str) -> dict:
    return {"chunks": [], "embeddings_ref": None}

def causal_propose(variables: list[str], priors: dict | None = None) -> dict:
    return {"dags": [], "scores": []}

def causal_intervene(dag: dict, do: dict) -> dict:
    return {"expected_delta": 0.0, "plan": []}

def evidence_verify(claim_id: str, source: str) -> dict:
    return {"verdict": "unknown", "score": 0.0}
```

6) Program-of-Thought Auditor

`engines/program_of_thought/ir.py`

```
from dataclasses import dataclass
@dataclass
class Step:
    name: str; inputs: dict; outputs: dict | None = None
@dataclass
class PoTIR:
    steps: list[Step]
```

adk_app/tools/smt_tools.py

```
from engines.program_of_thought.ir import PoTIR

def pot_compile(plan_text: str) -> dict:
    # TODO: parse to IR + invariants
    return {"ir": {"steps": []}, "invariants": []}

def smt_check(ir: dict) -> dict:
    # TODO: translate to SMT, run z3
    return {"status": "unknown", "model": None}

def proof_export(format: str = "lean") -> dict:
    return {"artifact": None}
```

adk_app/tools/fuzz_tools.py

```
from hypothesis import given, strategies as st

def fuzz_casegen(ir: dict, targets: list[str] | None = None) -> dict:
    # TODO: generate counterexamples for targets
    return {"counterexamples": []}
```

7) World-Simulation Planner

engines/world_sim/model.py

```
class WorldModel:
    def __init__(self):
        self.params = {}
    def predict(self, state, action):
        # TODO: learned transition + cost/latency predictor
        return state, {"cost": 0.0, "risk": 0.0}
```

adk_app/tools/world_model_tools.py

```
from engines.world_sim.model import WorldModel
WM = WorldModel()

def sim_train(history: list[dict]) -> dict:
    return {"ref": "wm:0"}
```

```

def sim_rollout(goals: dict, tools: list[str], budget: dict) -> dict:
    # TODO: beam search plans with EVC
    return {"ranked_plans": []}

def act_execute(plan: dict) -> dict:
    return {"events": [], "artifacts": [], "costs": {}}

```

adk_app/tools/guard_tools.py

```

def guard_estimate_risk(plan: dict) -> dict:
    return {"risk_score": 0.0, "red_flags": []}

def rollback_recover(snapshot_id: str) -> dict:
    return {"status": "restored"}

```

8) MCP Server (mcp_server/server.py)

```

import json, sys
from pathlib import Path
from fastapi import FastAPI
from pydantic import BaseModel

app = FastAPI()

# Load schemas dynamically
SCHEMA_DIR = Path(__file__).parent / "schemas"
REGISTRY_DIR = Path(__file__).parent / "registry"
REGISTRY_DIR.mkdir(parents=True, exist_ok=True)

@app.get("/tools")
def list_tools():
    tools = []
    for p in SCHEMA_DIR.glob("*.json"):
        tools.append(json.loads(p.read_text()))
    for p in REGISTRY_DIR.glob("*.json"):
        tools.append(json.loads(p.read_text()))
    return {"tools": tools}

@app.post("/call/{name}")
def call_tool(name: str, body: dict):
    # TODO: route to actual implementations (import modules, validate schemas)
    return {"name": name, "echo": body}

```

9) Security & Policy

- **Allowlist:** `adk_app/policies/tool_allowlist.yaml` defines callable tools; block shadow/dup names.
 - **Quotas:** CPU/mem/time/net limits enforced in `sandbox/policy.py` for generated tools.
 - **HIL gates:** Any egress/spend/write tool returns `requiresApproval: true` for the agent to request confirmation.
 - **Taint tags:** Mark externally fetched content; sanitize before reuse.
 - **Audit:** Log all tool calls with `{user, ts, tool, args_hash, outcome}`.
-

10) Demos

10.1 90-second Tool-Foundry Demo

```
$ ./scripts/run_demo.sh tool-foundry
# Prompts agent: "Extract tables from this scanned PDF and join with a CSV..."
# Expected: tool_specify → tool_codegen → tool_sandbox_run → tool_register →
call tool → output table JSON
```

10.2 Causal Card Demo

```
$ ./scripts/run_demo.sh causal
# Ingest 3 papers; output claims, DAG hypotheses, suggested interventions,
confidence
```

10.3 PoT Auditor Demo

```
$ ./scripts/run_demo.sh pot
# Compile reasoning to IR; z3 check; export proof; fuzz for counterexamples
```

10.4 World-Sim Planner Demo

```
$ ./scripts/run_demo.sh world-sim
# Train tiny world model from synthetic tool histories; simulate plans; execute
top plan with guardrails
```

11) Testing

```
pytest -q
pytest tests/test_tool_foundry.py::test_codegen_register
pytest tests/test_pot_auditor.py::test_smt_contracts
```

12) Next Steps (Fill-In Priority)

1. Implement `pdf_tables` OCR+table extraction with fallback chain.
 2. Replace spec/codegen heuristics with LLM-guided spec planner constrained by JSON schema.
 3. Add `tda_summarize` persistence diagrams (Giotto-tda or Ripser) and cluster mapping.
 4. Flesh out `pot_compile` translation templates and Z3 contracts for DP/graph problems.
 5. Add EVC in world-sim with cost/risk prediction and rollback snapshots.
 6. Harden sandbox: seccomp profile, file IO deny-list, absolute wall clock cap.
-

13) README.md (starter copy)

Shock Agent Suite delivers four novel agents: 1) Tool-Foundry (self-extending tools) 2) Chronosynaptic causal research 3) Program-of-Thought auditor 4) World-simulation planner.

Built for Google AI Agent Intensive. ADK function tools + MCP server. Sandbox tests, proofs, and cost/risk ledgers included.

```
make dev    # or scripts/dev_up.sh
make test
make demo TOOL=tool-foundry
```

License: Choose permissive (Apache-2.0) or keep private during the course.